

Exercise 2: Programming Distributed Systems (Summer 2025)


Submission

Deadline: 13.05.25 AOE

- You need a team and a Gitlab repository for this exercise sheet.
- In your Git repository, create a branch for this exercise sheet (for example with `git checkout -b ex2`)
- Create a folder named “ex2” in your repository and add your solutions to this folder.
- Test your submission with the provided test cases¹. Feel free to add more tests, but do not change the existing test cases.
- Make use of Dialyzer annotations to improve your code quality

1 Time and causality

- Give an example execution that shows that for two timestamps from Lamport clocks, $C(e_1) < C(e_2)$ does not imply that $e_1 \rightarrow e_2$.
- Given a configuration of three servers S^1, S^2, S^3 and their current corresponding vector-clocks $S_C^1 = [1, 3, 2]$ $S_C^2 = [1, 4, 2]$ $S_C^3 = [1, 0, 3]$. Is such a configuration possible? If yes, draw an example execution. If no, explain why.



$t(b_1) = 1 < 2 = t(a_2)$, but $b_1 \not\rightarrow a_2$.

Not possible. S_3 receives an event from S_1 , but S_1 also receives at least one event from S_3 . Locally, the vector clock increases on sending and on receiving events. Therefore, the first component of S_1 cannot be 1. The configuration is impossible.

2 Implementing Vector Clocks

A vector clock is a mapping from processes to positive integers². Implement a module named `VectorClock` with the following functions:

- `new()` creates a new vector clock, where all processes have value 0.
- `increment(vc, p)` increments the entry of process `p` by 1.

¹You can use `mix test` for executing the tests

²In the literature it is often assumed that processes are numbered which allows to write down clocks like $[4, 7, 3]$ or $\begin{pmatrix} 4 \\ 7 \\ 3 \end{pmatrix}$ instead of the longer $\{p_1 \mapsto 4, p_2 \mapsto 7, p_3 \mapsto 3\}$. However, in this exercise we do not assume that the number of processes is known and arbitrary terms can be used as process names.

- `get(vc, p)` returns the value for process `p`.
- `leq(vc1, vc2)` checks, whether `vc1` is less than or equal to `vc2`. This is the case, iff $\forall p. \text{get}(vc_1, p) \leq \text{get}(vc_2, p)$.
- `merge(vc1, vc2)` merges two vector clocks by computing their least upper bound (the smallest vector clock `v`, such that $vc_1 \leq v$ and $vc_2 \leq v$).

Feel free to add more tests, but do not change the existing test cases.

3 Testing Vector Clocks

Specify at least 3 invariants that should hold for your vector clock implementation. The tests should be using these invariants, derive at least 2 test cases for each invariant and for the following example invariant:

1. `new() == new()`

Example invariants:

- $\forall v, p: v \leq \text{increment}(v, p)$
 - Increment by one strictly increases size
- $\forall v: \text{new}() \leq v$
 - A new vectorclock is smaller or equal to any other vectorclock
- $\forall v, v_2: v \leq \text{merge}(v, v_2)$
 - Merging a vectorclock produces a bigger or equal vectorclock
- $\forall v: v \leq v$
 - \leq –reflexivity
- $\forall v_1, v_2, v_3: v_1 \leq v_2 \wedge v_2 \leq v_3 \rightarrow v_1 \leq v_3$
 - \leq –transitivity
- $\forall v_1, v_2: v_1 \leq v_2 \wedge v_2 \leq v_1 \rightarrow v_1 == v_2$
 - \leq –antisymmetry
- $\forall v_1, v_2, v_3: \text{merge}(\text{merge}(v_1, v_2), v_3) == \text{merge}(v_1, \text{merge}(v_2, v_3))$
 - Merge associativity
- $\forall v_1, v_2, \forall (p, i) \in \text{merge}(v_1, v_2) : (p, i) \in v_1 \vee (p, i) \in v_2$
 - Merge by components
- $\forall v_1, v_2. v' = \text{merge}(v_1, v_2) : v_1 \leq v_2 \rightarrow v' == v_2$
 - Merge subsumption
- $\forall v_1, v_2. \text{merge}(v_1, v_2) == \text{merge}(v_1, \text{merge}(v_1, v_2))$
 - Merge is idempotent

Be careful with using structural equality `==`! Sometimes, data structures have metadata attached that should be ignored in the equality check.